

Hadoop: Distributed Processing of Big Data

Lecture 01



About Me

- Marilson Campos
 - Principal Data Architect with 15+ years of experience.
 - Background in Machine Learning and Search engines
 - Designing Big Data systems since 2007
- Some relevant presentations
 - Hadoop Summit 2013
 - Apache Impala User Group - 2015
 - Strata 2016 – Use case @ Cloudera Booth
 - Impala spotlight – Cloudera 2016

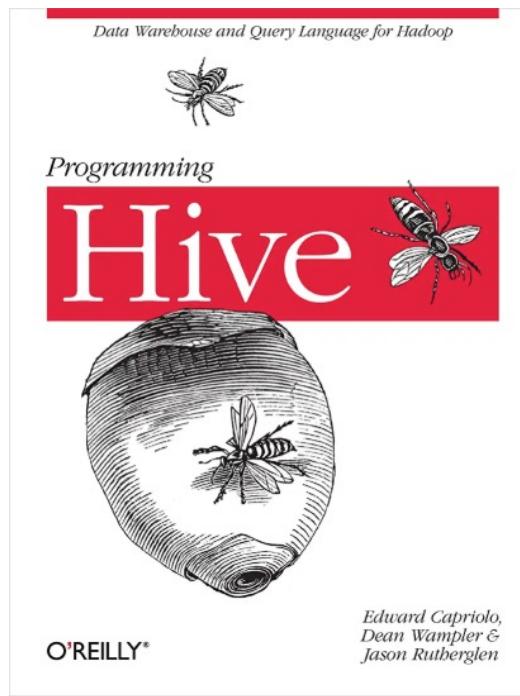
About the class

What we will cover

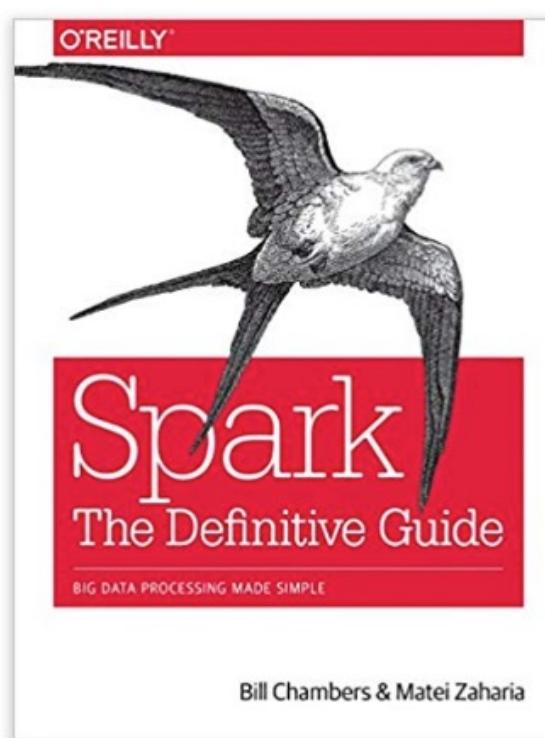
- Big Data architecture and its components
- Hadoop Map/Reduce model
- Spark programming.
- Big Data using Hive (SQL)
- Extending Hive
- Kafka & HBase
- Data pipeline design

Recommended Books

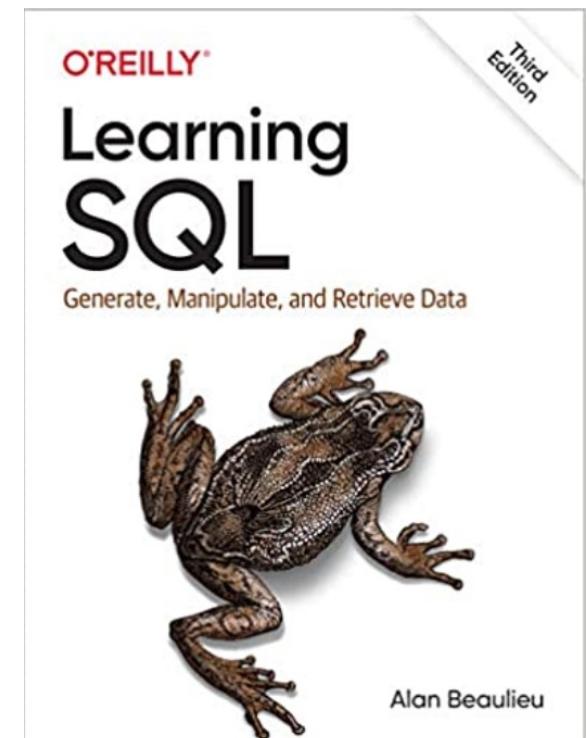
Hive



Spark



SQL



About the class

At the end of this class, I expect you to...

- Have an intuition about Big Data system design.
- Be able to write Big data programs.
- Know how to design a “data pipeline”.
- Continue to learn using other sources.

About the class

Class Material vs The Test

- We are covering some non-trivial concepts.
 - If you miss few advanced details, don't worry too much.
Important: If you are lost, please stop me! And ask for clarification.
 - I want to make the class as interactive as possible.
- Problems in the labs are a lot harder than the ones in the test.
 - This is by design, so advanced students can also benefit from the class.
- To succeed in the test, you need to master the foundational knowledge about a Hadoop cluster, Hive, and Spark.

Tests

Test 1 – After 5th Lecture (Take home)

- 8 questions worth 5 points each.
 - 7 multiple choice questions
 - 1 question with one paragraph answer.

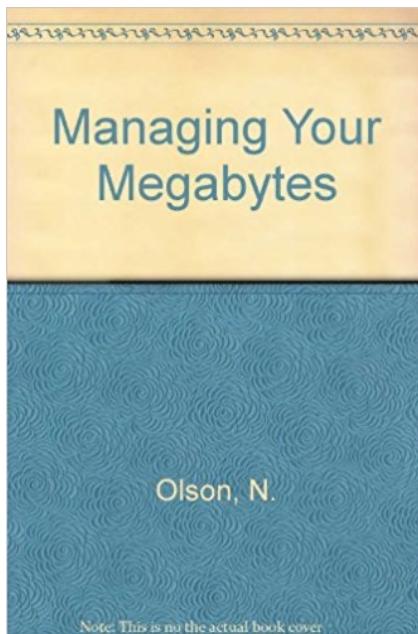
Test 2 – At 8th Lecture - Open book.

- 4 questions worth 15 points each.
 - Test your knowledge on perform operations like:
 - Load file into a Hadoop cluster
 - Create a table mapping to files loaded.
 - Perform transformations on datasets.
 - Query and extract results from cluster.

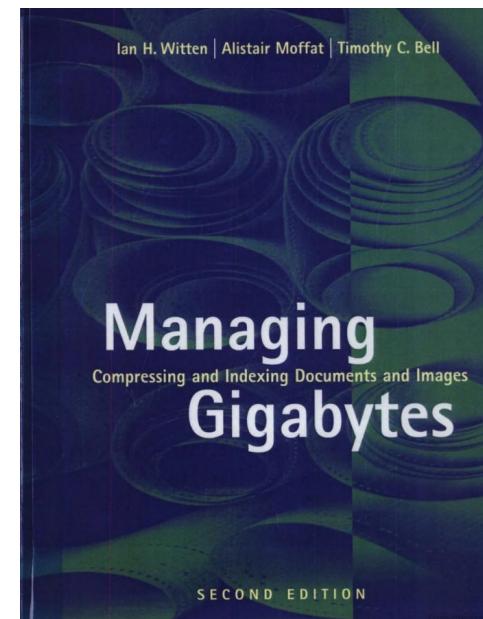
Big Data

Why Big Data?

- More information being stored digitally.
- Need to process larger datasets and unstructured data.
- Cannot process in traditional databases systems.



Published 1986



Published 1994

A New Architecture is needed

Data Transfer versus Storage Size

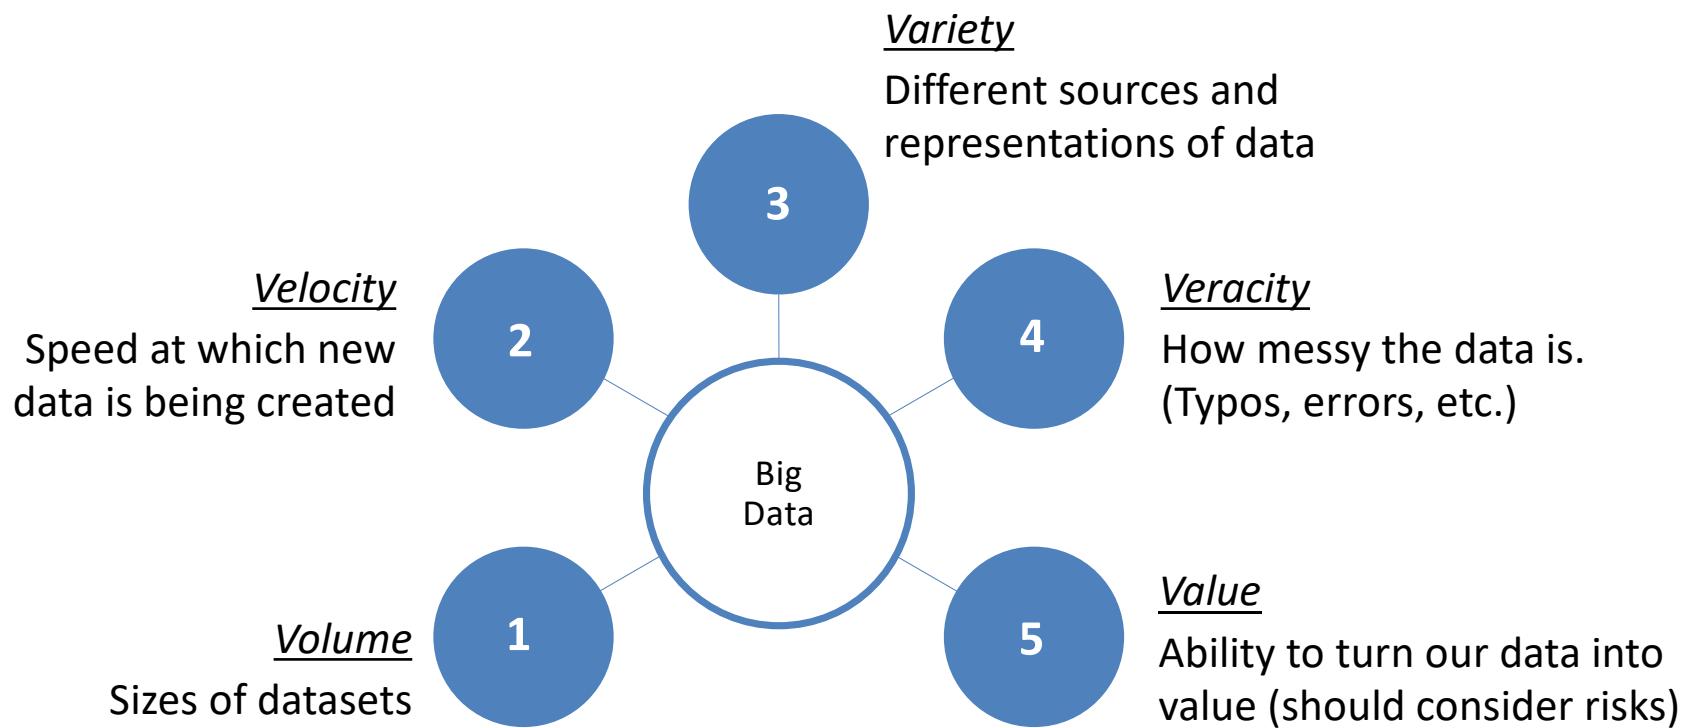
- 1990
 - 1Gb Storage
 - 4Mb per second data transfer
 - Time to read all data = **approx. 4 minutes**
- 2019
 - 4,000Gb Storage (~4Tb)
 - 100Mb per second data transfer
 - Time to read all data = **approx. 11 hours.**

1Mb data transfer

Operation	X Speed factor
Processor Read	1000 x
Memory Read	120 x
SSD Read	20 x
Disk Read	1
Disk Seek	$\frac{1}{2}$ x
1Gpbs Net	$\frac{1}{2}$ x

Big Data

The 5 V's of Big Data



A New Architecture is Needed
High level abstraction

- Inverse of Virtualization
 - Virtualization = one machine appears to be many.
 - Big Data = many machines appear to be one.

A New Architecture is Needed Forces

- Represent large datasets
 - Distributed
 - Memory, disk or combination.
- Process datasets
 - Use cluster resources to execute in parallel.
- Handle failure
 - Automatically recover from failure.

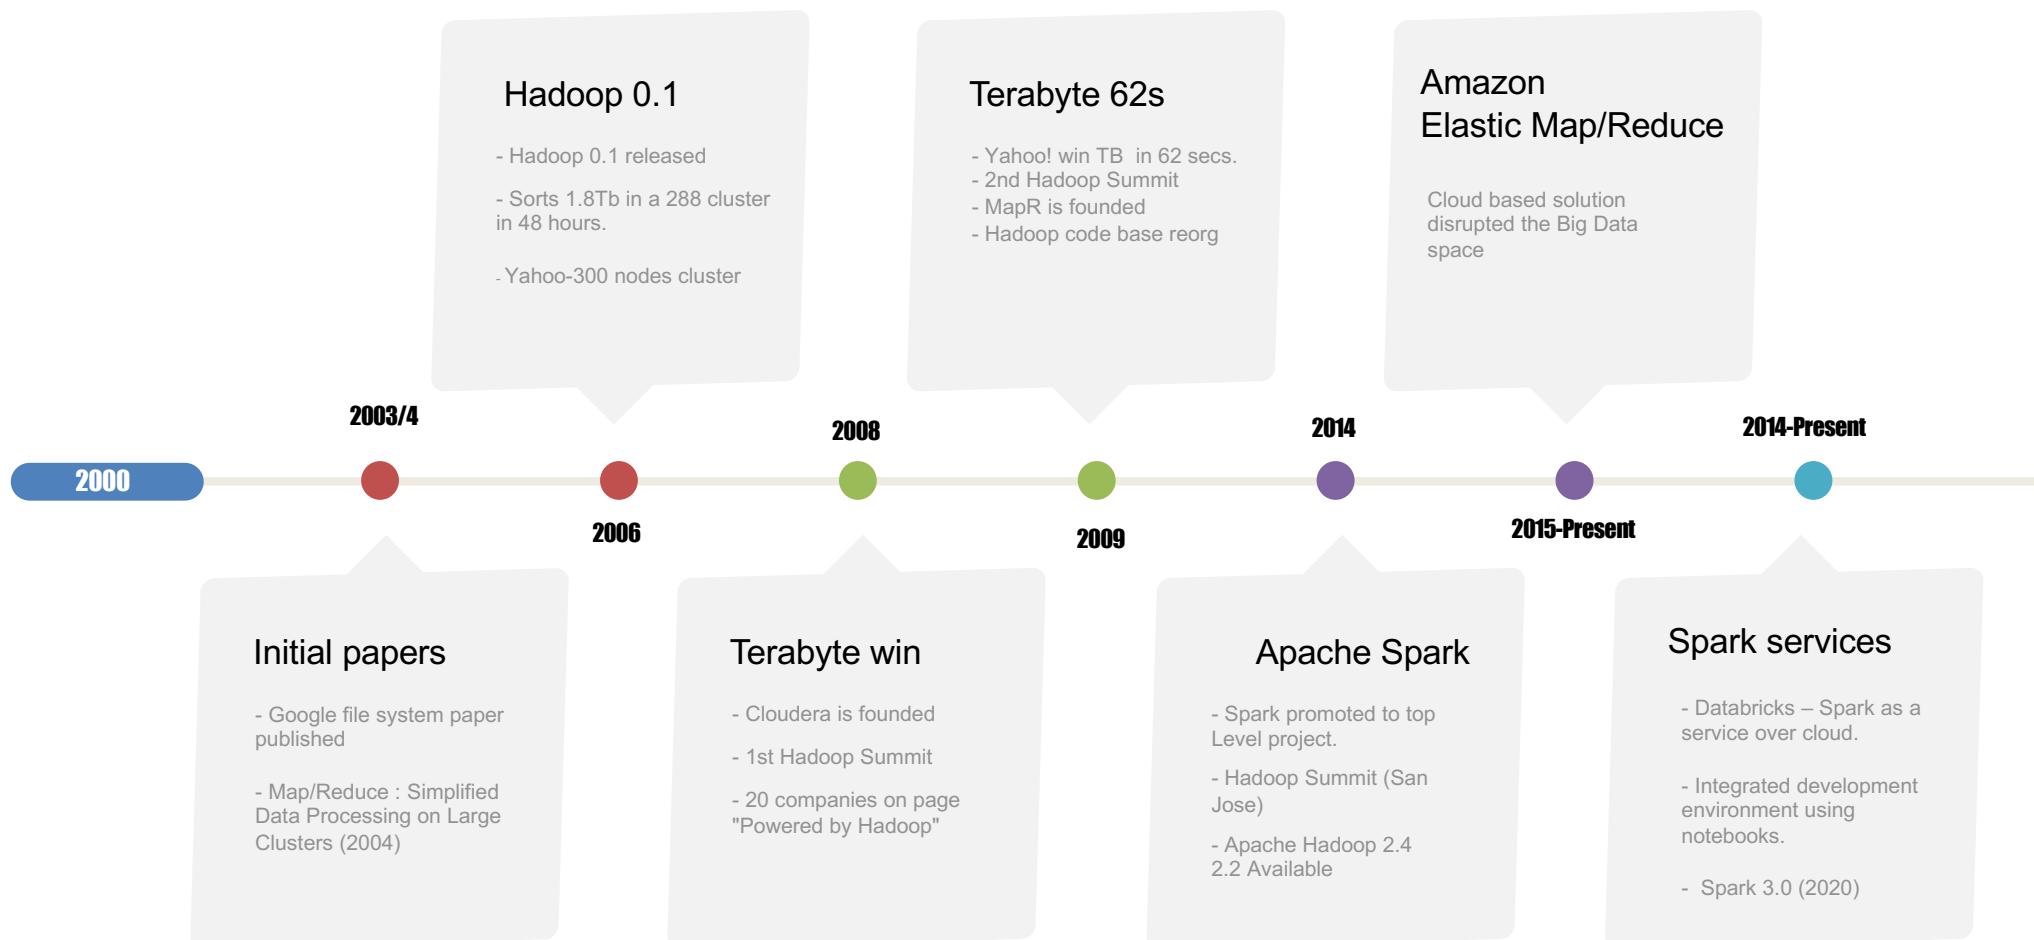
A New Architecture is Needed

Initial approach

- Big Data distributed System
 - Many Computers processing in parallel.
 - Many disks available.
 - No transaction beyond files.
 - Single computers/disks should be able to fail with no impact.

Hadoop is born

The Hadoop Timeline



Hadoop is born Big Data Distributions

- **In premises (Datacenter)**
 - **Cloudera**
 - Most popular in premises
 - Hybrid cloud offerings
 - **Hortonworks**
 - Most portable (Compatible with Windows)
 - Merged with Cloudera
 - **MapR**
 - Based on custom Linux filesystem
 - Acquired by HP.
- **Cloud Based**
 - **Amazon Elastic Map/Reduce (EMR)**
 - Cloud based simple to use
 - Current cloud-based leader.
 - **Microsoft**
 - Azure Big data
 - Part of Microsoft cloud offerings.
 - **Google**
 - Google compute engine
 - Part of Google other cloud offerings.

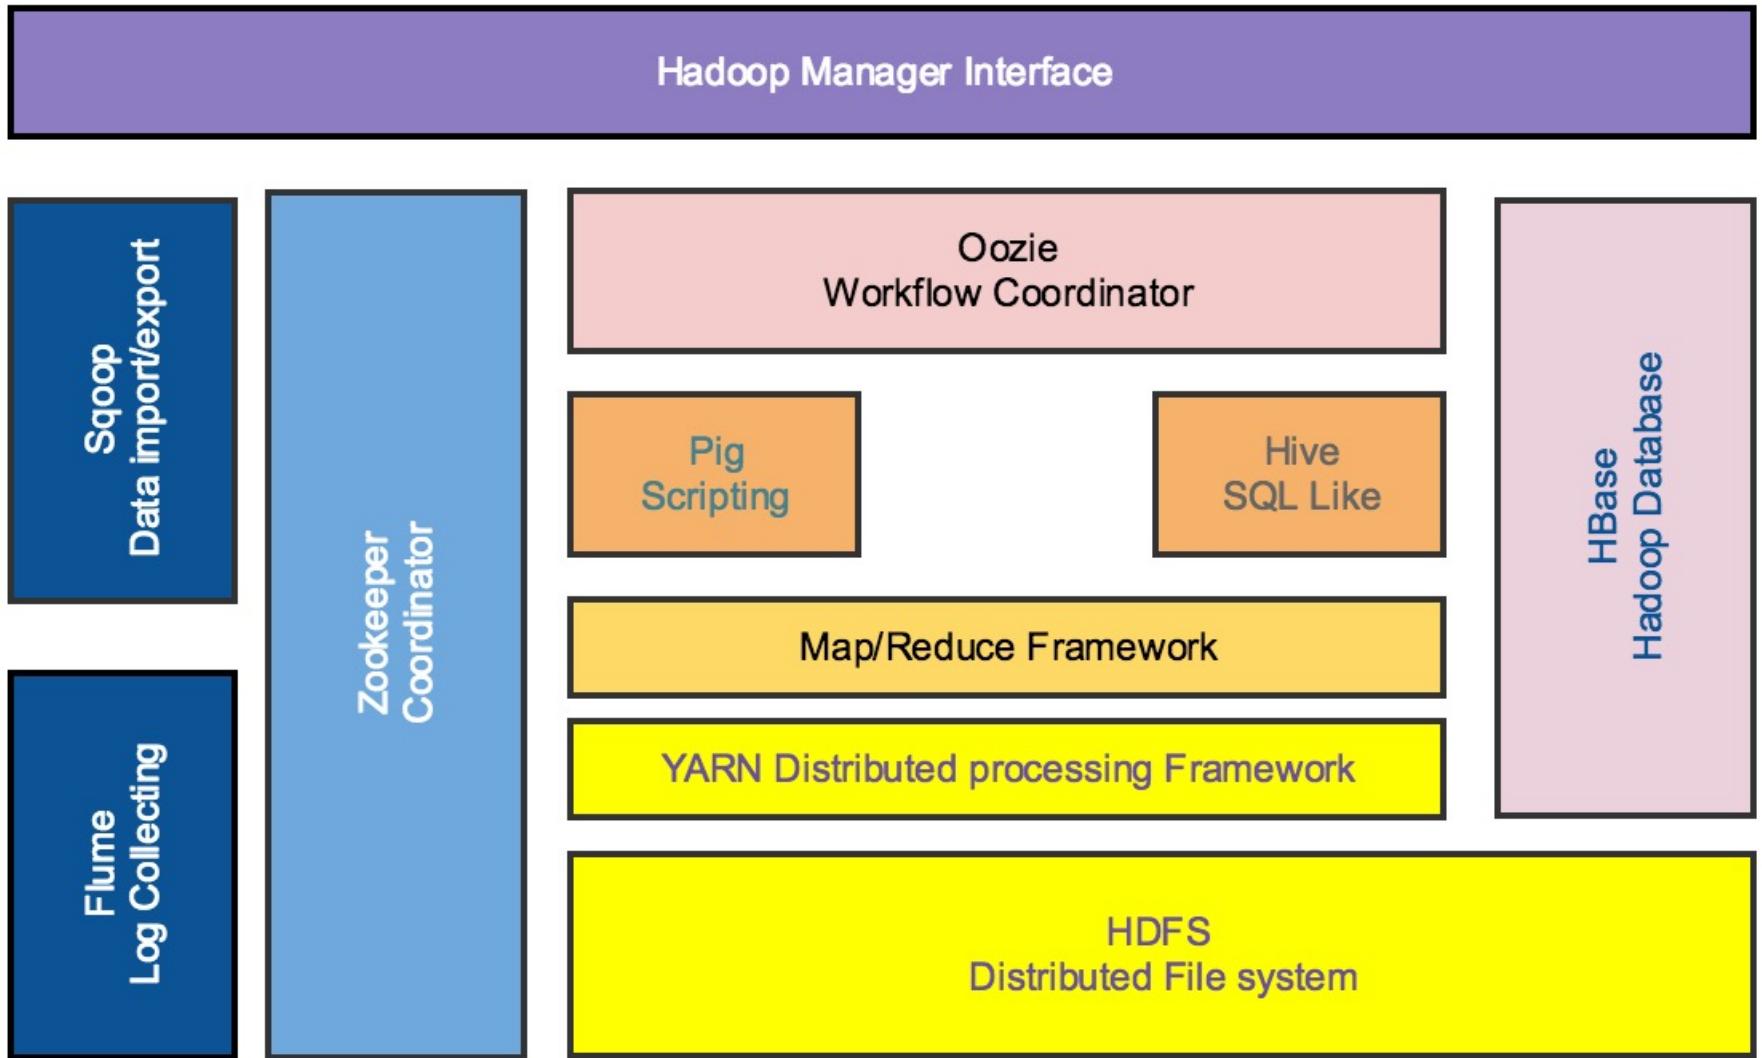
Hadoop is born

The Hadoop Cluster

- 1000's of servers running
 - Each with dozens of disks.
 - High chance of single failure.
- Automated recovery
 - In case of a failure, the system needs to detect it and recover.
 - Cluster programs should not fail in this case.

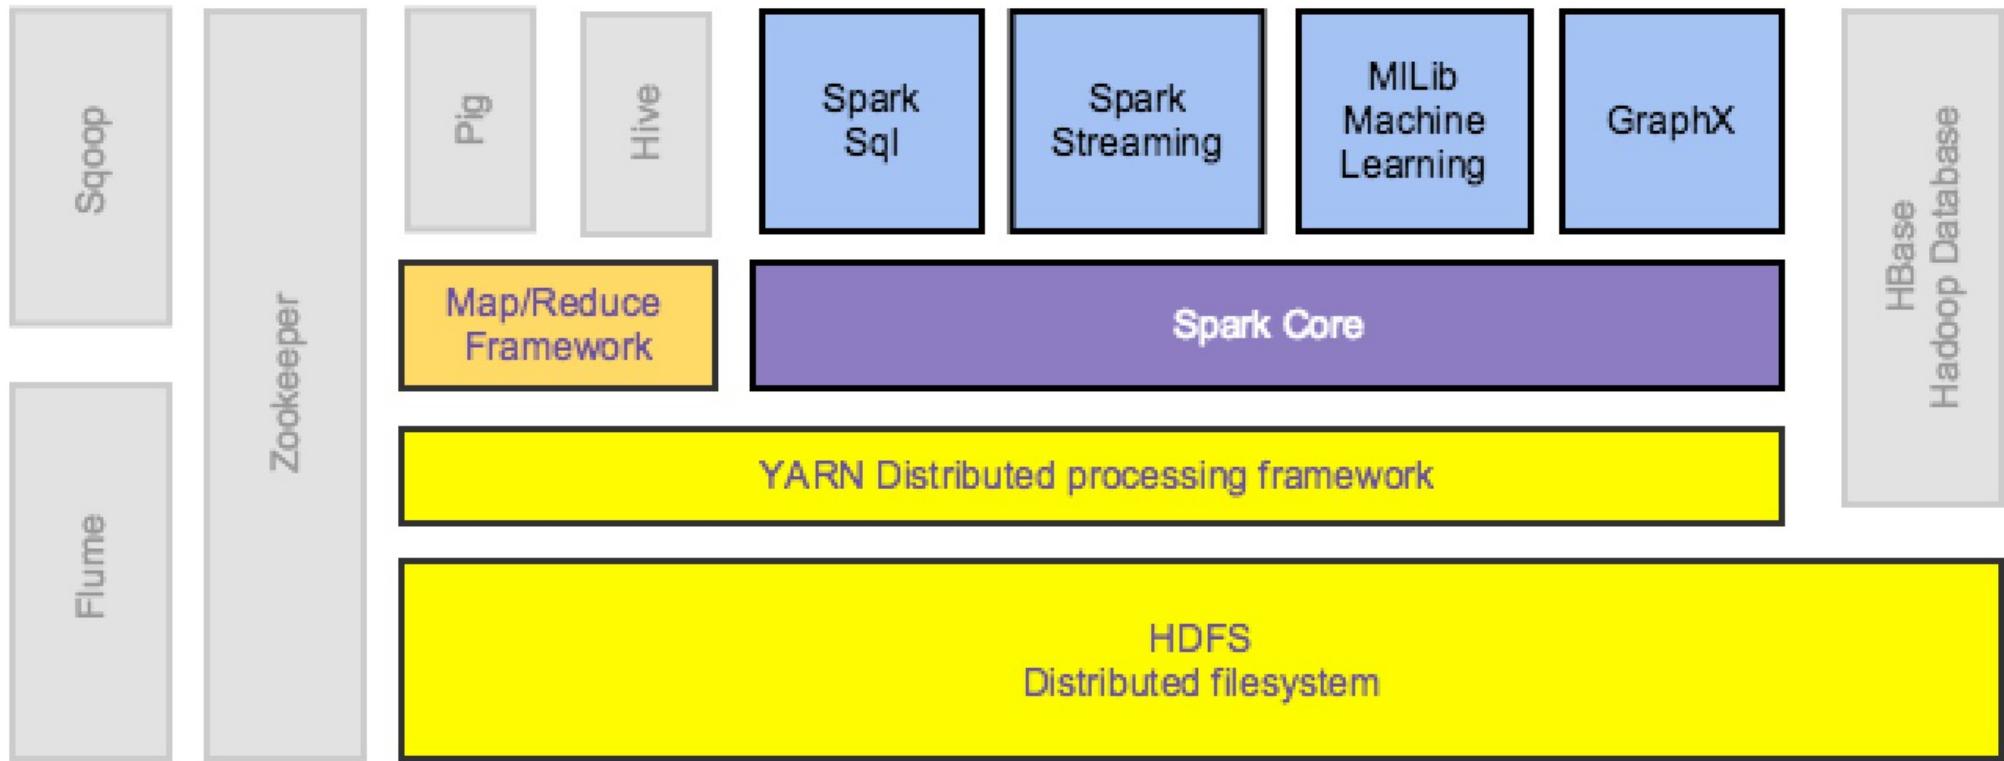
Hadoop is born

The Hadoop Stack



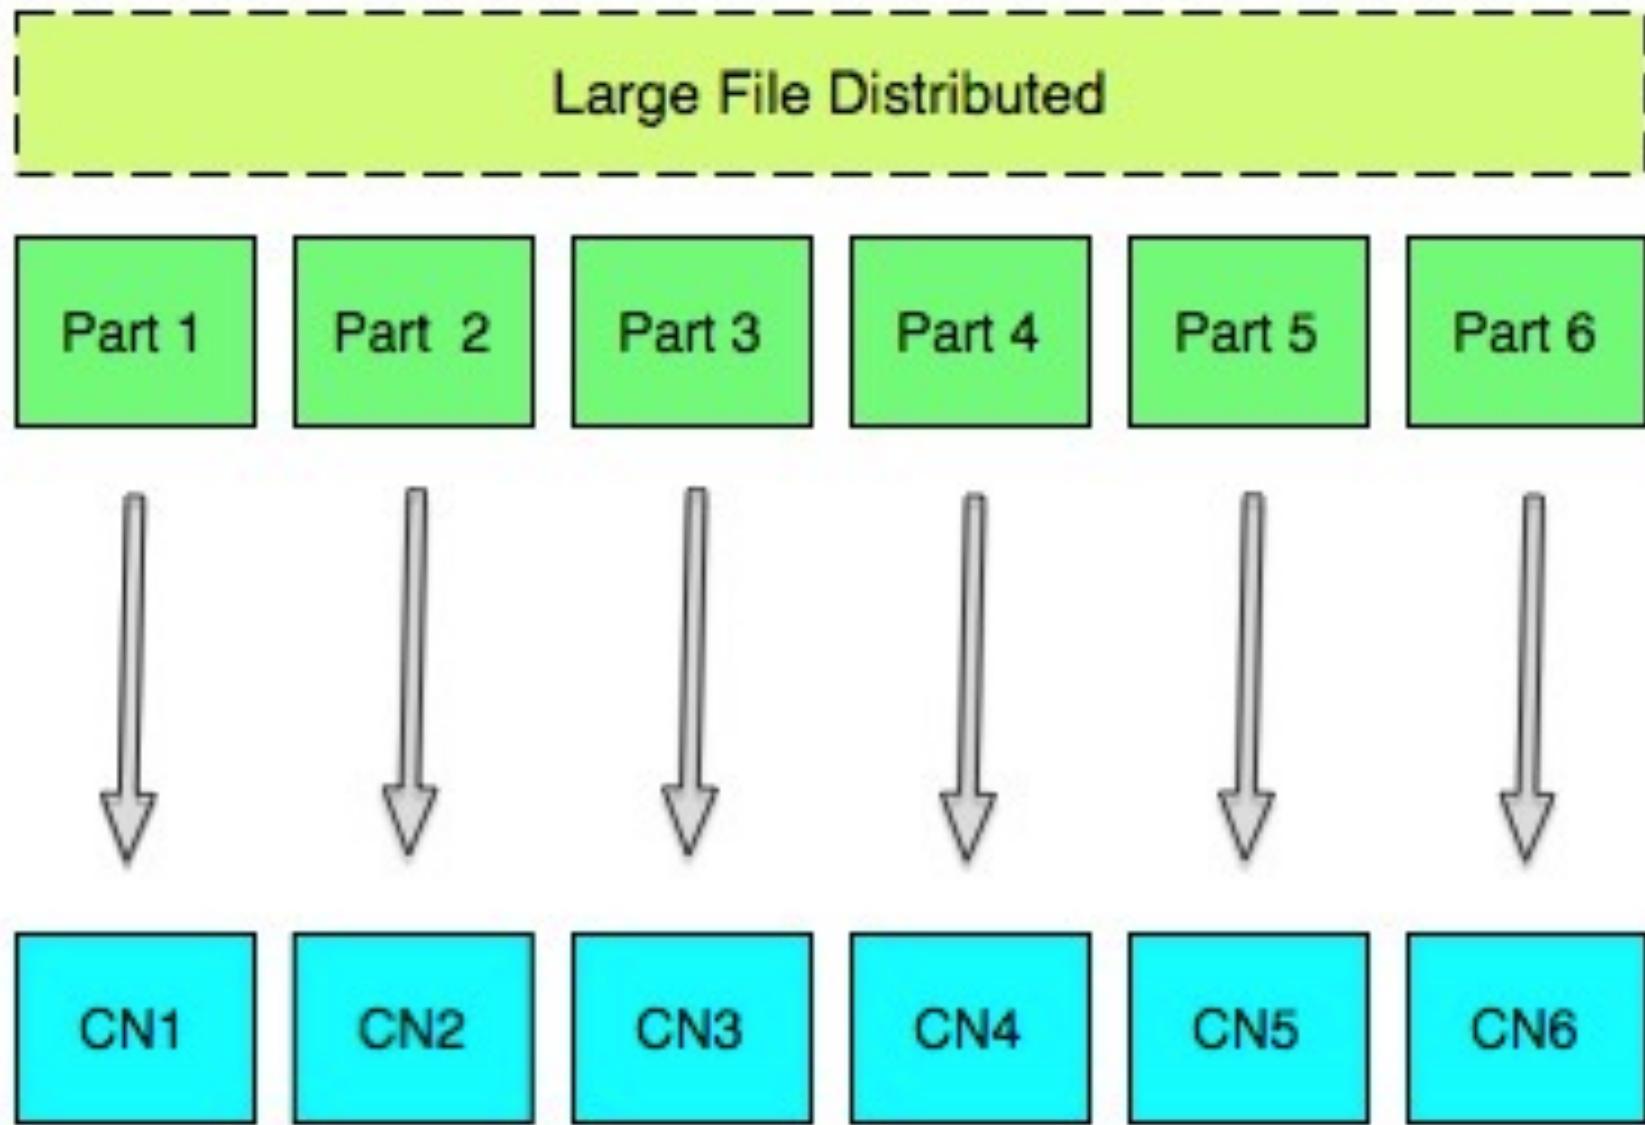
Hadoop is born

The Hadoop/Spark Stack



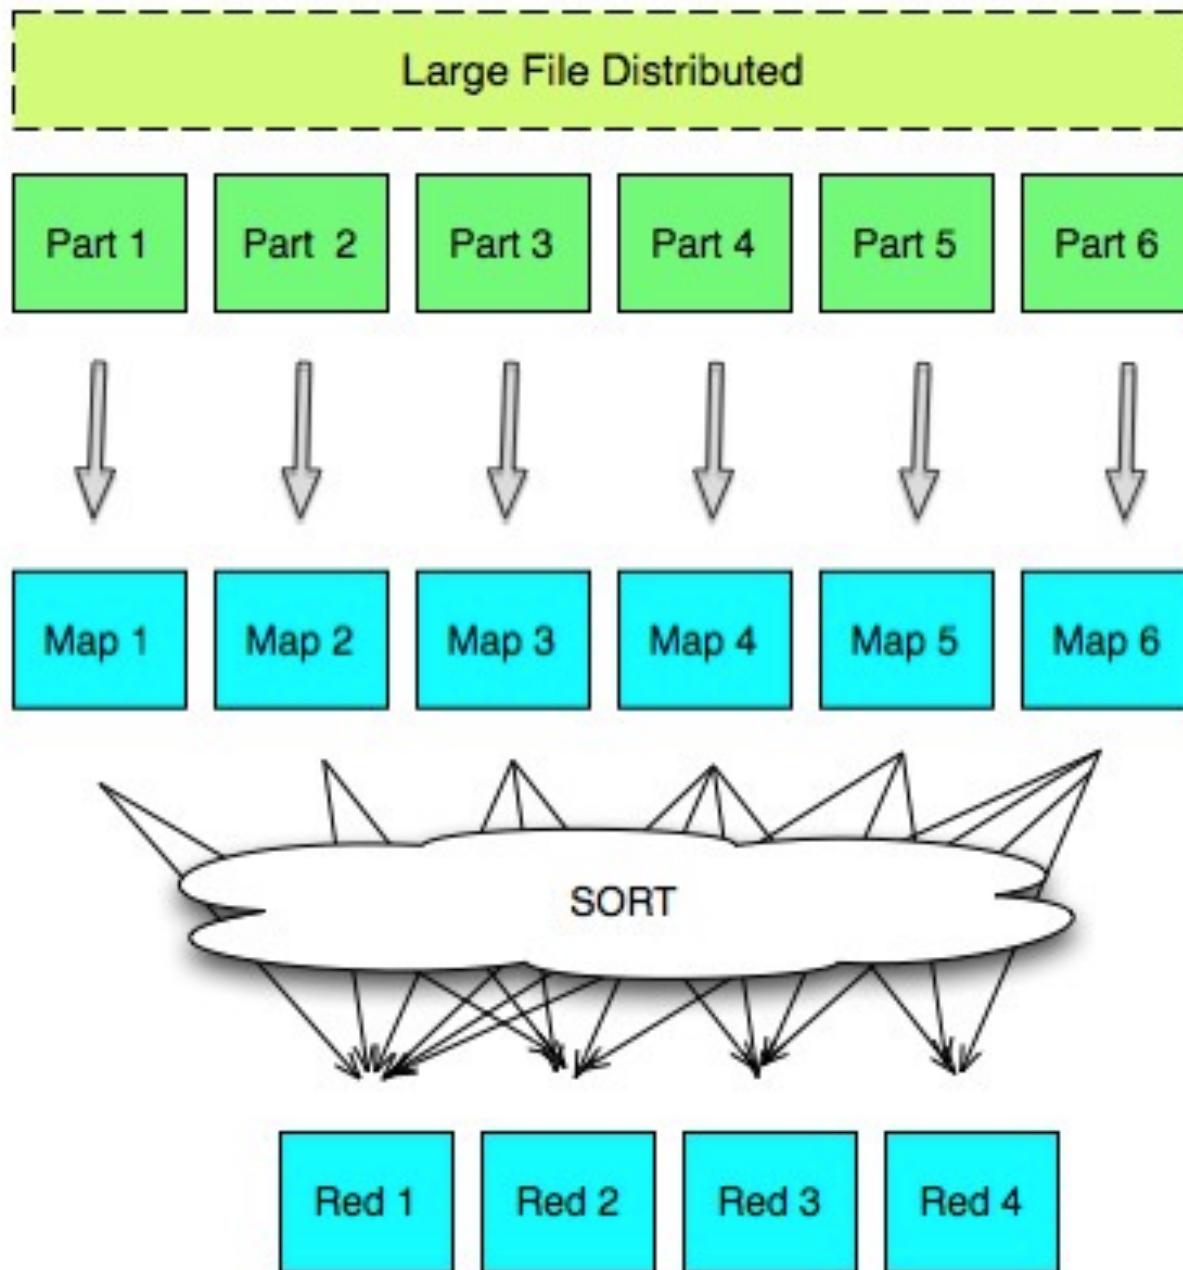
- Spark uses Hadoop core to execute jobs.
- Spark uses HDFS for storage.

MapReduce
Distributed execution v0.1 – pure map



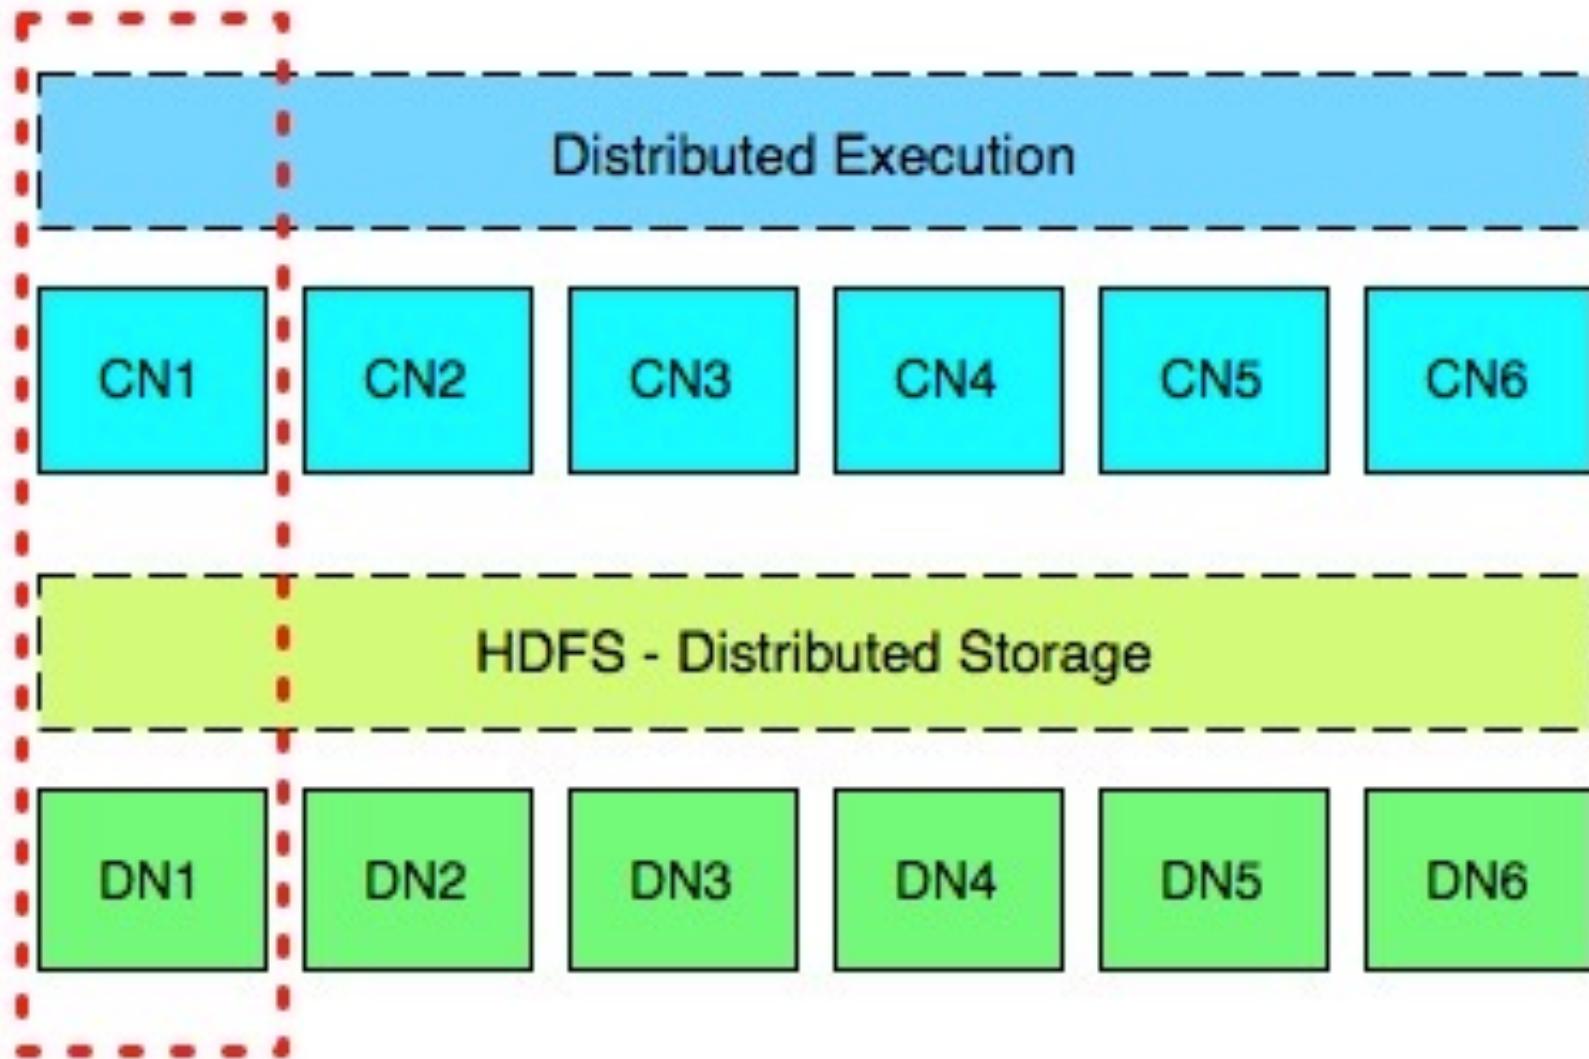
MapReduce

Distributed execution v1 – map/reduce



MapReduce

Distributed execution – deployment



Hadoop is born

Hadoop Map/Reduce x Spark

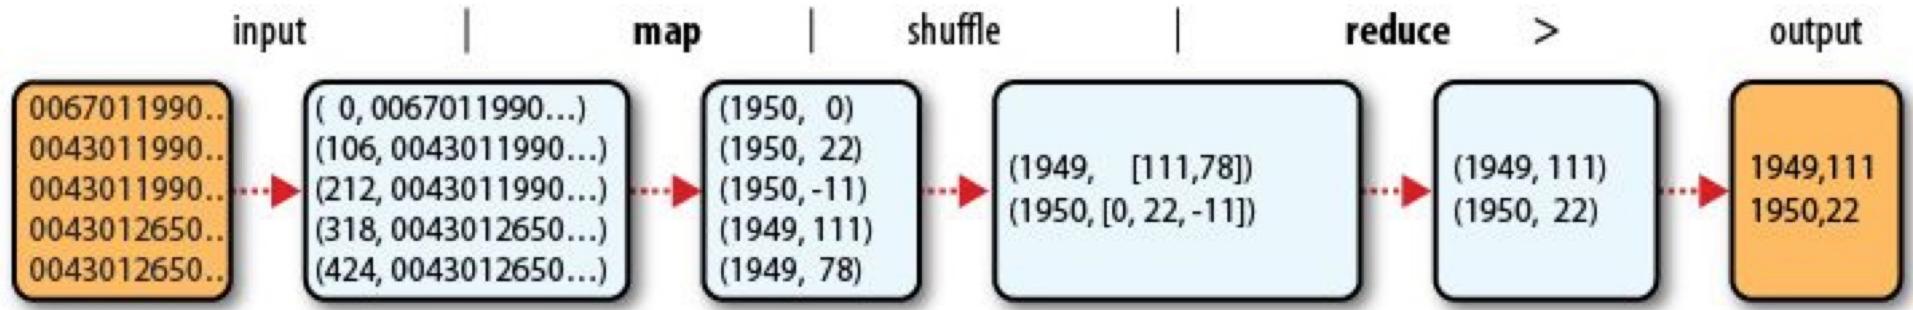


- Low level API
- Designed for batch processing (slower starting time)
- Long running process
- Reads and writes data to and from disk many times.



- Higher level API
- Tries to use and keep data in memory.
- Supports batch and near real-time.
(faster starting time)

Hadoop Map/Reduce Map/Reduce processing



Data flow:

- Input
- Map (transformation)
- Shuffle (sort)
- Reduce
- Output

Hadoop Map/Reduce

The Mapper

```
public class MaxTemperatureMapper
    extends Mapper<LongWritable, Text, Text, IntWritable> {

    private static final int MISSING = 9999;

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String line = value.toString();
        String year = line.substring(15, 19);
        int airTemperature;
        if (line.charAt(87) == '+') { // parseInt doesn't like leading plus signs
            airTemperature = Integer.parseInt(line.substring(88, 92));
        } else {
            airTemperature = Integer.parseInt(line.substring(87, 92));
        }
        String quality = line.substring(92, 93);
        if (airTemperature != MISSING && quality.matches("[01459]")) {
            context.write(new Text(year), new IntWritable(airTemperature));
        }
    }
}
```

Hadoop Map/Reduce

The Reducer

```
import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class MaxTemperatureReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context)
        throws IOException, InterruptedException {

        int maxValue = Integer.MIN_VALUE;
        for (IntWritable value : values) {
            maxValue = Math.max(maxValue, value.get());
        }
        context.write(key, new IntWritable(maxValue));
    }
}
```

Hadoop Map/Reduce

The Driver

```
public static void main(String[] args) throws Exception {  
    if (args.length != 2) {  
        System.err.println("Usage: MaxTemperature <input path> <output path>");  
        System.exit(-1);  
    }  
  
    Job job = new Job();  
    job.setJarByClass(MaxTemperature.class);  
    job.setJobName("Max temperature");  
  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.setMapperClass(MaxTemperatureMapper.class);  
    job.setReducerClass(MaxTemperatureReducer.class);  
  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

Hadoop Map/Reduce - Run output 1/2

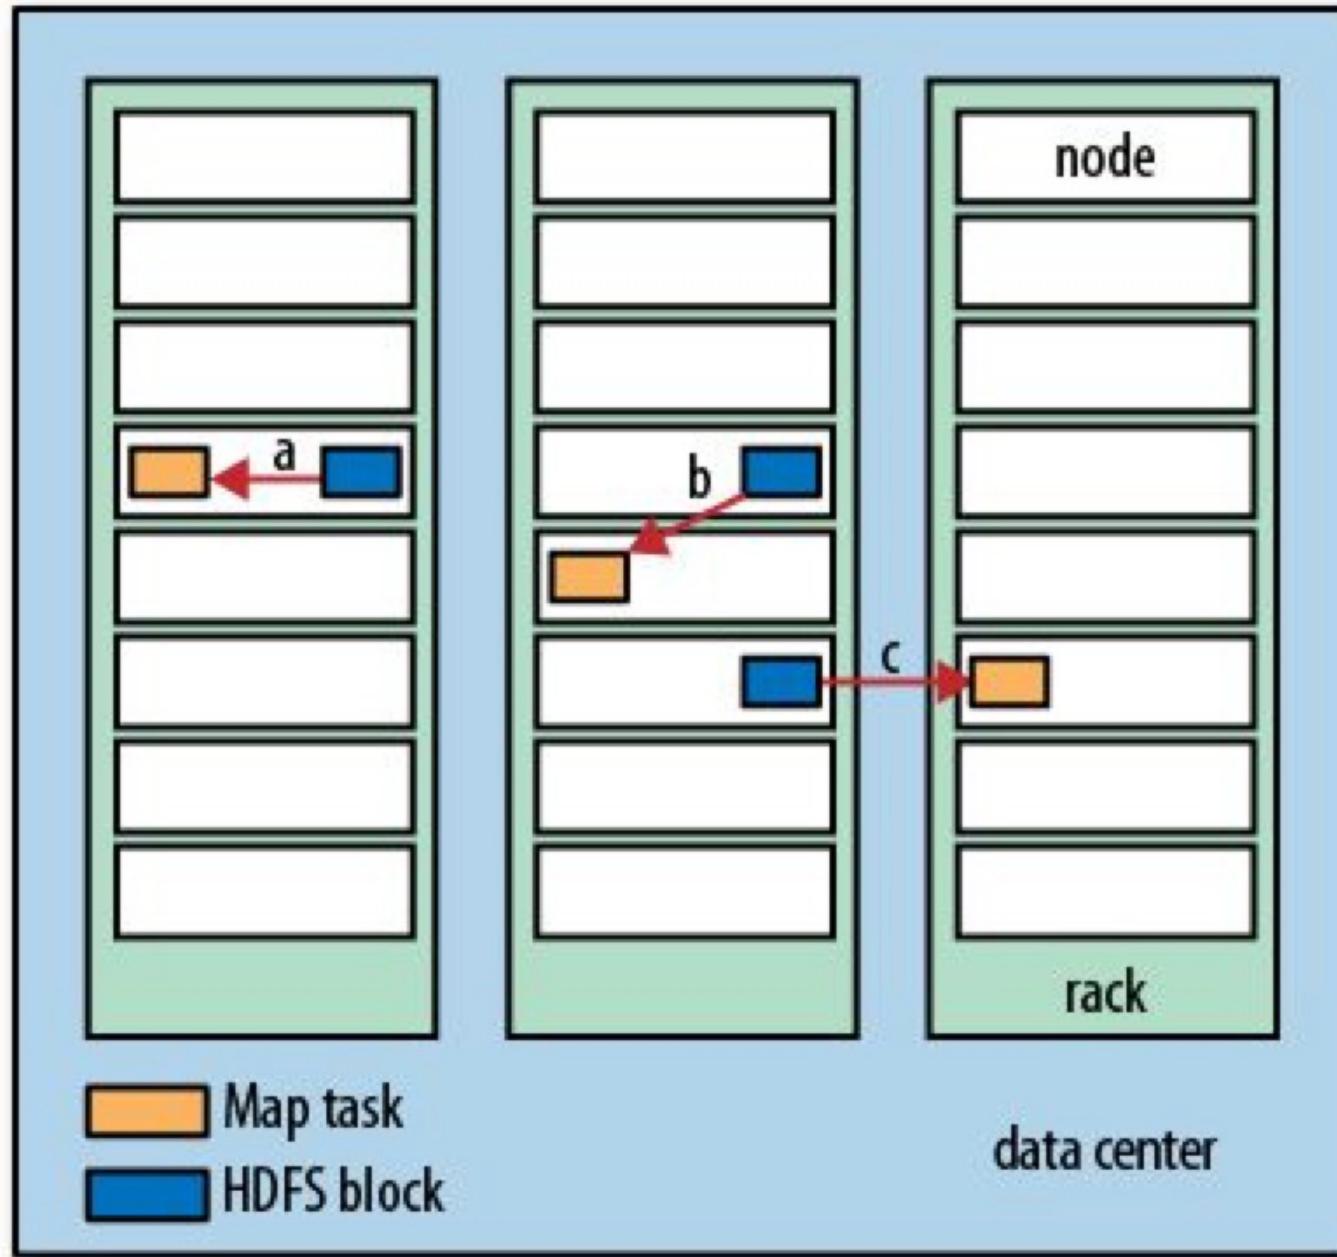
```
12/02/04 11:50:41 WARN util.NativeCodeLoader: Unable to load native-hadoop library  
for your platform... using builtin-java classes where applicable  
12/02/04 11:50:41 WARN mapred.JobClient: Use GenericOptionsParser for parsing the  
arguments. Applications should implement Tool for the same.  
12/02/04 11:50:41 INFO input.FileInputFormat: Total input paths to process : 1  
12/02/04 11:50:41 INFO mapred.JobClient: Running job: job_local_0001  
12/02/04 11:50:41 INFO mapred.Task: Using ResourceCalculatorPlugin : null  
12/02/04 11:50:41 INFO mapred.MapTask: io.sort.mb = 100  
12/02/04 11:50:42 INFO mapred.MapTask: data buffer = 79691776/99614720  
12/02/04 11:50:42 INFO mapred.MapTask: record buffer = 262144/327680  
12/02/04 11:50:42 INFO mapred.MapTask: Starting flush of map output  
12/02/04 11:50:42 INFO mapred.MapTask: Finished spill 0  
12/02/04 11:50:42 INFO mapred.Task: Task:attempt_local_0001_m_000000_0 is done. And i  
s in the process of committing  
12/02/04 11:50:42 INFO mapred.JobClient: map 0% reduce 0%  
12/02/04 11:50:44 INFO mapred.LocalJobRunner:  
12/02/04 11:50:44 INFO mapred.Task: Task 'attempt_local_0001_m_000000_0' done.  
12/02/04 11:50:44 INFO mapred.Task: Using ResourceCalculatorPlugin : null  
12/02/04 11:50:44 INFO mapred.LocalJobRunner:  
12/02/04 11:50:44 INFO mapred.Merger: Merging 1 sorted segments  
12/02/04 11:50:44 INFO mapred.Merger: Down to the last merge-pass, with 1 segments  
left of total size: 57 bytes  
12/02/04 11:50:44 INFO mapred.LocalJobRunner:  
12/02/04 11:50:45 INFO mapred.Task: Task:attempt_local_0001_r_000000_0 is done. And
```

Hadoop Map/Reduce - Run output 2/2

```
is in the process of committing
12/02/04 11:50:45 INFO mapred.LocalJobRunner:
12/02/04 11:50:45 INFO mapred.Task: Task attempt_local_0001_r_000000_0 is allowed to
commit now
12/02/04 11:50:45 INFO output.FileOutputCommitter: Saved output of task 'attempt_local
_0001_r_000000_0' to output
12/02/04 11:50:45 INFO mapred.JobClient: map 100% reduce 0%
12/02/04 11:50:47 INFO mapred.LocalJobRunner: reduce > reduce
12/02/04 11:50:47 INFO mapred.Task: Task 'attempt_local_0001_r_000000_0' done.
12/02/04 11:50:48 INFO mapred.JobClient: map 100% reduce 100%
12/02/04 11:50:48 INFO mapred.JobClient: Job complete: job_local_0001
12/02/04 11:50:48 INFO mapred.JobClient: Counters: 17
12/02/04 11:50:48 INFO mapred.JobClient: File Output Format Counters
12/02/04 11:50:48 INFO mapred.JobClient: Bytes Written=29
12/02/04 11:50:48 INFO mapred.JobClient: FileSystemCounters
12/02/04 11:50:48 INFO mapred.JobClient: FILE_BYTES_READ=357503
12/02/04 11:50:48 INFO mapred.JobClient: FILE_BYTES_WRITTEN=425817
12/02/04 11:50:48 INFO mapred.JobClient: File Input Format Counters
12/02/04 11:50:48 INFO mapred.JobClient: Bytes Read=529
12/02/04 11:50:48 INFO mapred.JobClient: Map-Reduce Framework
12/02/04 11:50:48 INFO mapred.JobClient: Map output materialized bytes=61
12/02/04 11:50:48 INFO mapred.JobClient: Map input records=5
12/02/04 11:50:48 INFO mapred.JobClient: Reduce shuffle bytes=0
12/02/04 11:50:48 INFO mapred.JobClient: Spilled Records=10
12/02/04 11:50:48 INFO mapred.JobClient: Map output bytes=45
12/02/04 11:50:48 INFO mapred.JobClient: Total committed heap usage (bytes)=36923
8016
12/02/04 11:50:48 INFO mapred.JobClient: SPLIT_RAW_BYTES=129
12/02/04 11:50:48 INFO mapred.JobClient: Combine input records=0
12/02/04 11:50:48 INFO mapred.JobClient: Reduce input records=5
12/02/04 11:50:48 INFO mapred.JobClient: Reduce input groups=2
12/02/04 11:50:48 INFO mapred.JobClient: Combine output records=0
12/02/04 11:50:48 INFO mapred.JobClient: Reduce output records=2
12/02/04 11:50:48 INFO mapred.JobClient: Map output records=5
```

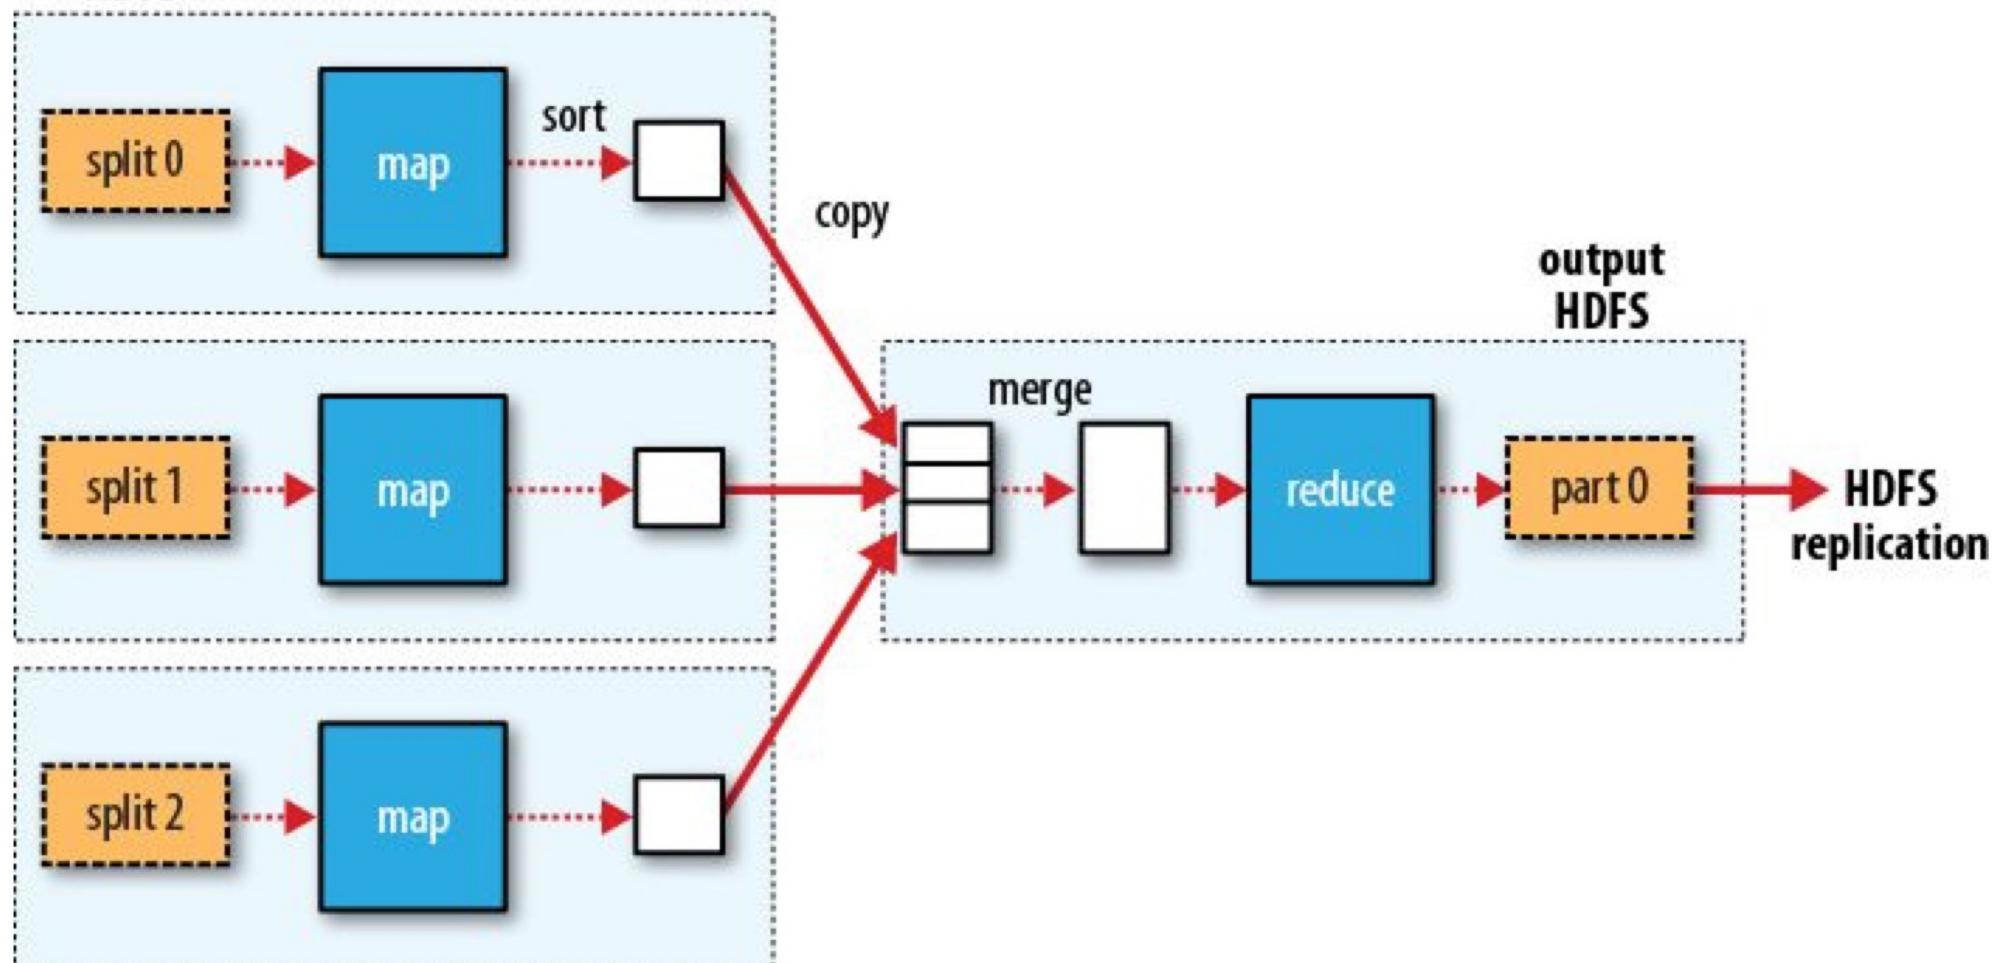
Hadoop Map/Reduce

Send the code to the data!



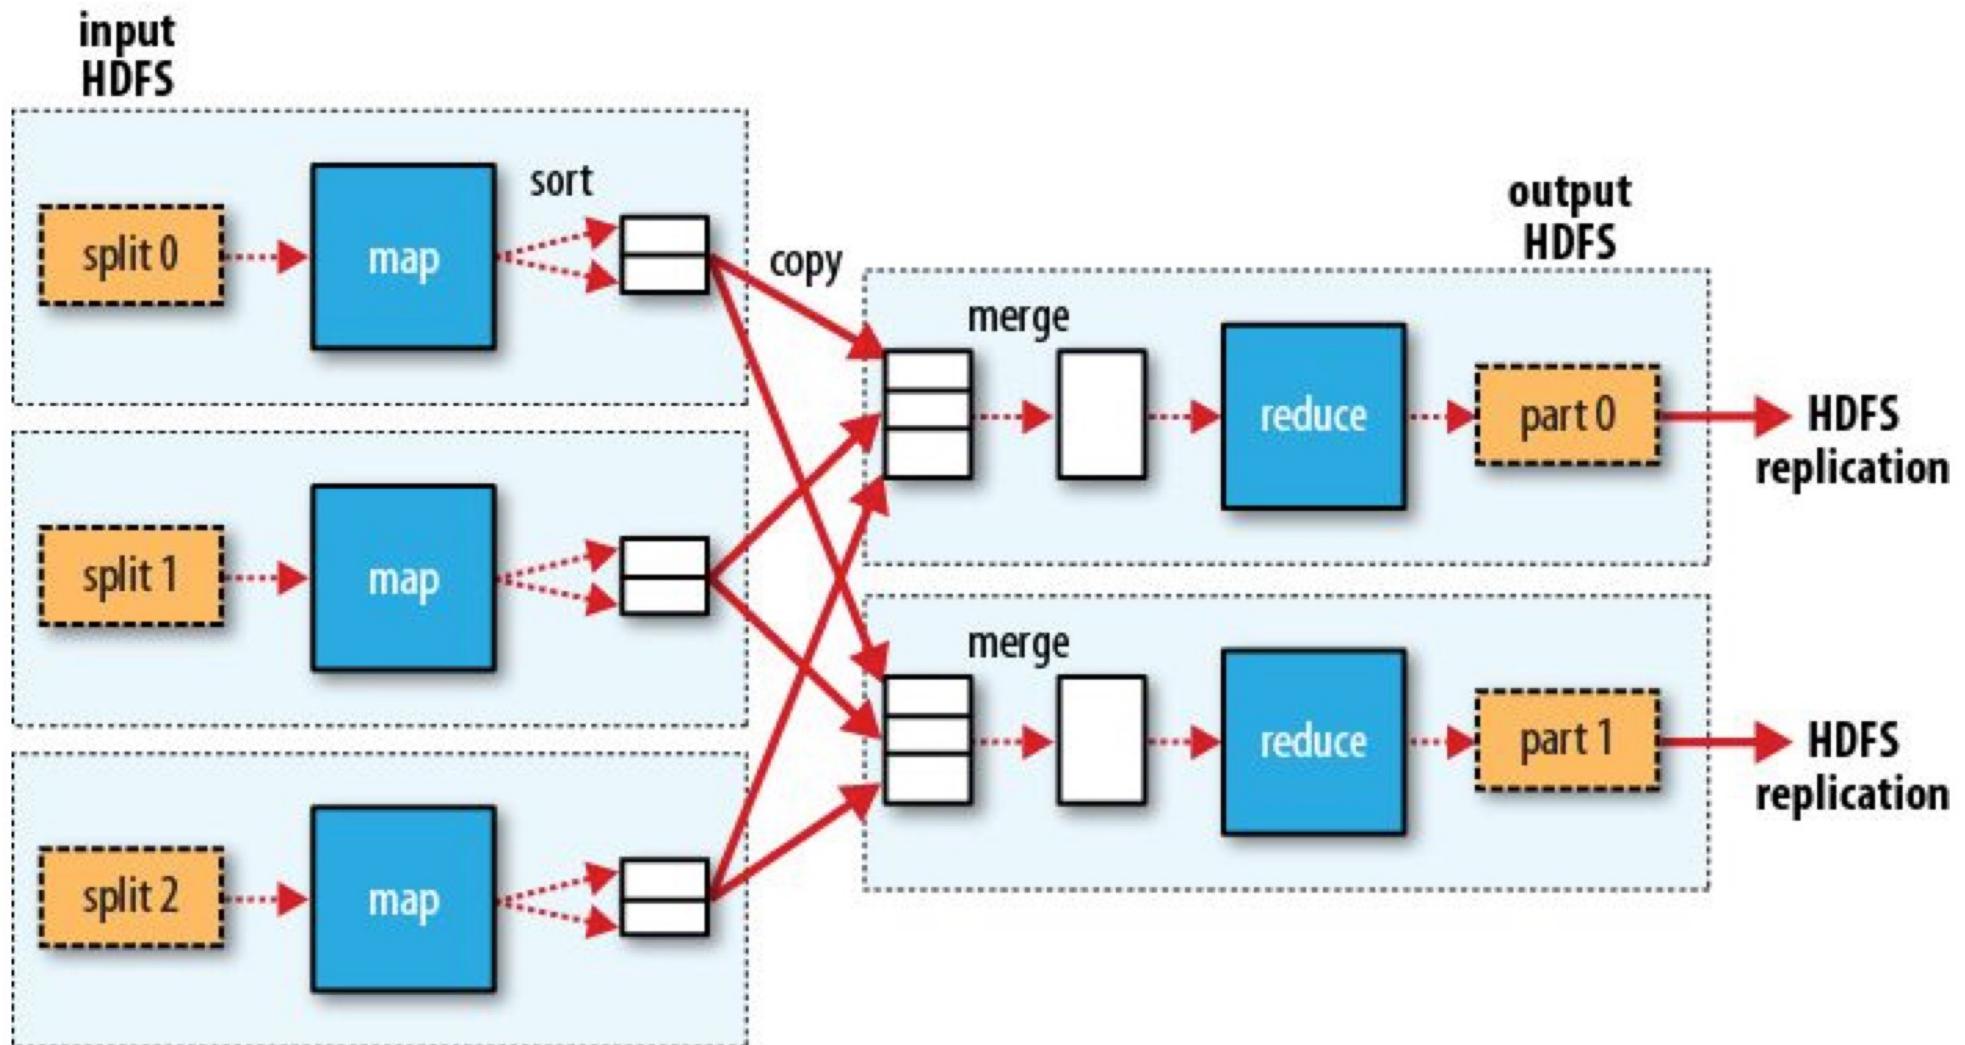
Hadoop Map/Reduce

MR Splits and dataflow 1 reducer



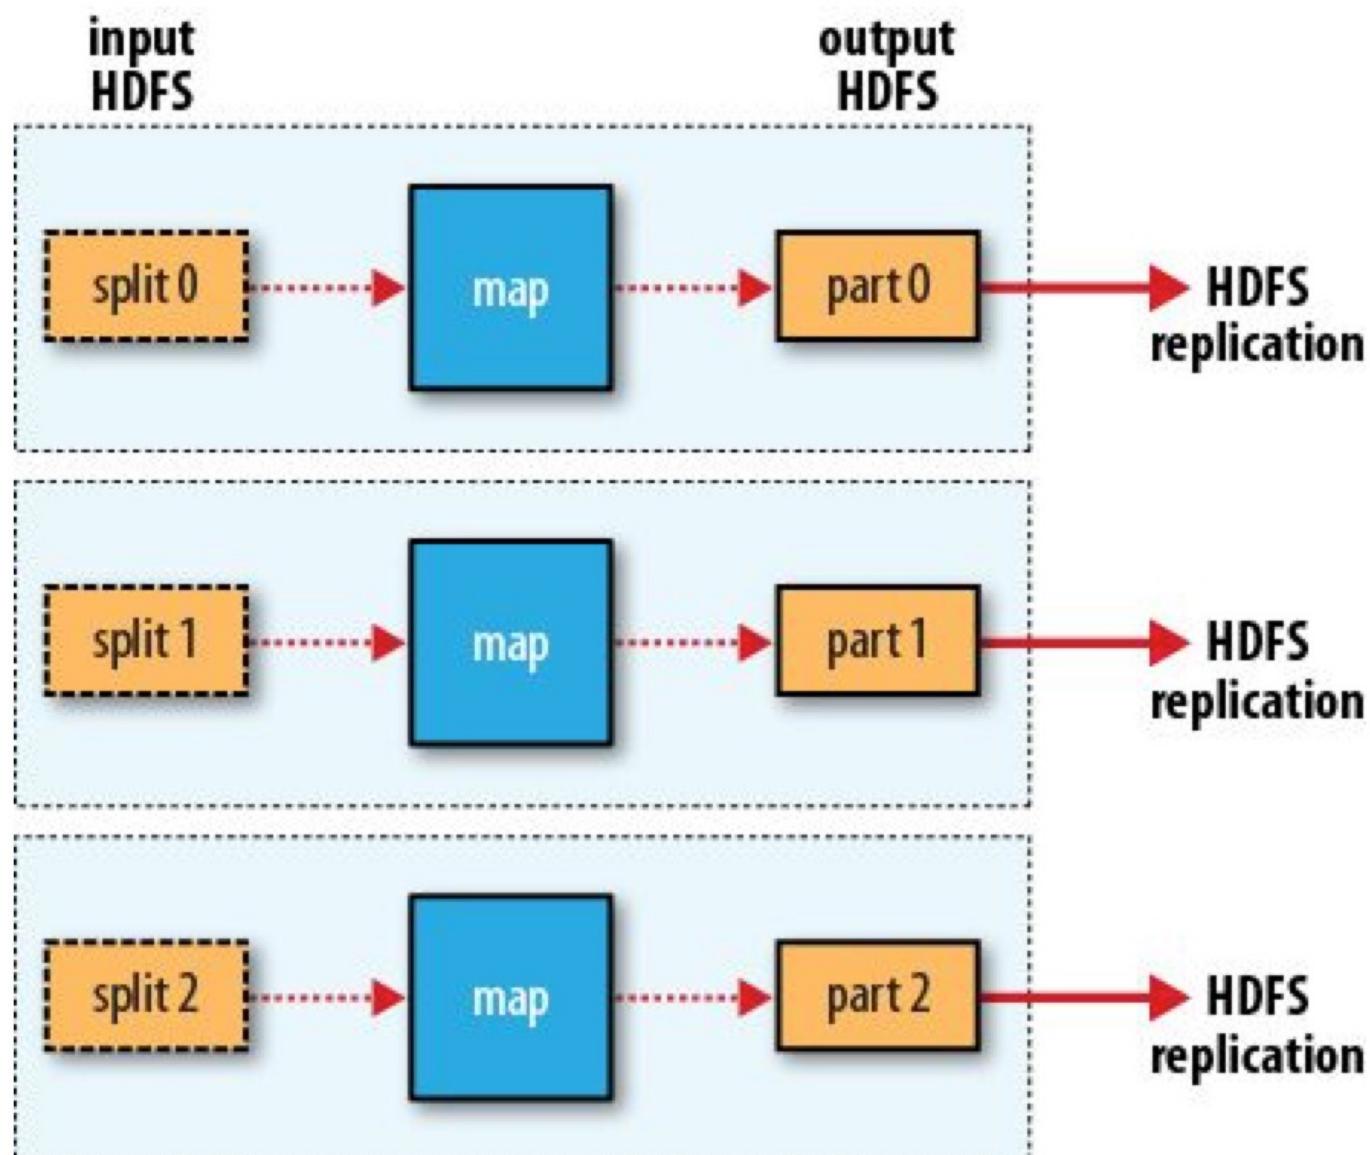
Hadoop Map/Reduce

MR Splits and dataflow multiple reducers



Hadoop Map/Reduce

MR Splits and dataflow Zero reducers



Hadoop Map/Reduce

The Combiner

- Calculates Intermediary results
- No guarantees if it will run or even how many times it will runs.
- Optimization feature

Hadoop Map/Reduce

The Driver with Combiner

```
public class MaxTemperatureWithCombiner {

    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Usage: MaxTemperatureWithCombiner <input path> " +
                "<output path>");
            System.exit(-1);
        }

        Job job = new Job();
        job.setJarByClass(MaxTemperatureWithCombiner.class);
        job.setJobName("Max temperature");

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setMapperClass(MaxTemperatureMapper.class);
        job.setCombinerClass(MaxTemperatureReducer.class);
        job.setReducerClass(MaxTemperatureReducer.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

Hadoop Map/Reduce Different integrations

- Map/Reduce Api -> JVM language (Java, Scala)
- Hadoop M/R Streaming -> Unix Stdin/Stdout.
- Hadoop pipes -> Unix Sockets

HDFS – Hadoop File Systems

Big Data Storage

- Unix like filesystem
- Access rights identical to Linux
- No support to updates
- Self healing capabilities.

HDFS – Hadoop File Systems

Big Data Storage

- Very large file support
- Optimized for continuous "reads".
- High throughput at the expense of latency
- Single writer, no support to multiple writers.

HDFS – Hadoop File Systems

HDFS Concepts

- **Blocks**
 - Atomic unit of storage
- **Namenode**
 - the master daemon for storage
 - Stores the metadata
 - 1 Active
 - 1 Standby
- **Journal Nodes**
 - Record changes on filesystem
 - Responsible for metadata updates on standby namenode.
- **Datanodes**
 - Responsive for reading/writing data.

HDFS – Hadoop File Systems

HDFS – Big Data Storage

- Namenodes
- Journal Nodes.
- Datanodes

Spark Spark lab Preparation

Signup for Databricks

Search for ‘Databricks Community Edition: Login’ and signup for the community edition.

Note: Do not create a trial account as it will expire and will require for you to pay for it.

Q & A